

Report

August 12, 2019

1 Continuous Control

In this notebook, you will learn how to use the Unity ML-Agents environment for the second project of the [Deep Reinforcement Learning Nanodegree](#) program.

1.0.1 1. Start the Environment

We begin by importing the necessary packages. If the code cell below returns an error, please revisit the project instructions to double-check that you have installed [Unity ML-Agents](#) and [NumPy](#).

```
[1]: from unityagents import UnityEnvironment
import numpy as np
```

Next, we will start the environment! *Before running the code cell below*, change the `file_name` parameter to match the location of the Unity environment that you downloaded.

- **Mac**: "path/to/Reacher.app"
- **Windows (x86)**: "path/to/Reacher_Windows_x86/Reacher.exe"
- **Windows (x86_64)**: "path/to/Reacher_Windows_x86_64/Reacher.exe"
- **Linux (x86)**: "path/to/Reacher_Linux/Reacher.x86"
- **Linux (x86_64)**: "path/to/Reacher_Linux/Reacher.x86_64"
- **Linux (x86, headless)**: "path/to/Reacher_Linux_NoVis/Reacher.x86"
- **Linux (x86_64, headless)**: "path/to/Reacher_Linux_NoVis/Reacher.x86_64"

For instance, if you are using a Mac, then you downloaded `Reacher.app`. If this file is in the same folder as the notebook, then the line below should appear as follows:

```
env = UnityEnvironment(file_name="Reacher.app")
```

```
[2]: env = UnityEnvironment(file_name='Reacher_Windows_x86_64/Reacher.exe')
```

```
INFO:unityagents:
'Academy' started successfully!
Unity Academy name: Academy
    Number of Brains: 1
```

```

    Number of External Brains : 1
    Lesson number : 0
    Reset Parameters :
        goal_speed -> 1.0
        goal_size -> 5.0
Unity brain name: ReacherBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
    Vector Observation space size (per agent): 33
    Number of stacked Vector Observation: 1
    Vector Action space type: continuous
    Vector Action space size (per agent): 4
    Vector Action descriptions: , , ,

```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```

[3]: # get the default brain
    brain_name = env.brain_names[0]
    brain = env.brains[brain_name]

```

1.0.2 2. Examine the State and Action Spaces

In this environment, a double-jointed arm can move to target locations. A reward of +0.1 is provided for each step that the agent's hand is in the goal location. Thus, the goal of your agent is to maintain its position at the target location for as many time steps as possible.

The observation space consists of 33 variables corresponding to position, rotation, velocity, and angular velocities of the arm. Each action is a vector with four numbers, corresponding to torque applicable to two joints. Every entry in the action vector must be a number between -1 and 1.

Run the code cell below to print some information about the environment.

```

[4]: # reset the environment
    env_info = env.reset(train_mode=True)[brain_name]

    # number of agents
    num_agents = len(env_info.agents)
    print('Number of agents:', num_agents)

    # size of each action
    action_size = brain.vector_action_space_size
    print('Size of each action:', action_size)

    # examine the state space
    states = env_info.vector_observations
    state_size = states.shape[1]
    print('There are {} agents. Each observes a state with length: {}'.
        →format(states.shape[0], state_size))

```

```
print('The state for the first agent looks like:', states[0])
```

Number of agents: 20

Size of each action: 4

There are 20 agents. Each observes a state with length: 33

The state for the first agent looks like: [0.00000000e+00 -4.00000000e+00

```
0.00000000e+00 1.00000000e+00
-0.00000000e+00 -0.00000000e+00 -4.37113883e-08 0.00000000e+00
0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
0.00000000e+00 0.00000000e+00 -1.00000000e+01 0.00000000e+00
1.00000000e+00 -0.00000000e+00 -0.00000000e+00 -4.37113883e-08
0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
0.00000000e+00 0.00000000e+00 5.75471878e+00 -1.00000000e+00
5.55726624e+00 0.00000000e+00 1.00000000e+00 0.00000000e+00
-1.68164849e-01]
```

1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Once this cell is executed, you will watch the agent's performance, if it selects an action at random with each time step. A window should pop up that allows you to observe the agent, as it moves through the environment.

Of course, as part of the project, you'll have to change the code so that the agent is able to use its experience to gradually choose better actions when interacting with the environment!

```
[6]: env_info = env.reset(train_mode=False)[brain_name]      # reset the environment
    →
states = env_info.vector_observations                        # get the current state
    → (for each agent)
scores = np.zeros(num_agents)                             # initialize the score
    → (for each agent)
while True:
    actions = np.random.randn(num_agents, action_size) # select an action (for
    → each agent)
    actions = np.clip(actions, -1, 1)                     # all actions between -1
    → and 1
    env_info = env.step(actions)[brain_name]              # send all actions to
    → the environment
    next_states = env_info.vector_observations             # get next state (for
    → each agent)
    rewards = env_info.rewards                            # get reward (for each
    → agent)
    dones = env_info.local_done                           # see if episode
    → finished
    scores += env_info.rewards                             # update the score (for
    → each agent)
```

```

    states = next_states                                # roll over states to
    →next time step
    if np.any(dones):                                  # exit loop if episode
    →finished
        break
print('Total score (averaged over agents) this episode: {}'.format(np.
    →mean(scores)))

```

Total score (averaged over agents) this episode: 0.07699999827891588

When finished, you can close the environment.

```
[ ]: env.close()
```

1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

1.0.5 Implementation details

`experiences = random.choices(self.memory, k=self.batch_size)` - randomness with replacement is a key for success of DDPG implementation.

4.1 Importing libraries

```
[1]: import datetime
import time
import random
import copy
from collections import namedtuple, deque
from unityagents import UnityEnvironment
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

```

4.2 Starting environment

```
[2]: env = UnityEnvironment(file_name='Reacher_Windows_x86_64/Reacher.exe')
```

```
INFO:unityagents:
'Academy' started successfully!
Unity Academy name: Academy
    Number of Brains: 1

```

```

    Number of External Brains : 1
    Lesson number : 0
    Reset Parameters :
        goal_speed -> 1.0
        goal_size -> 5.0
Unity brain name: ReacherBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
    Vector Observation space size (per agent): 33
    Number of stacked Vector Observation: 1
    Vector Action space type: continuous
    Vector Action space size (per agent): 4
    Vector Action descriptions: , , ,

```

4.3 Setting parameters

```

[3]: BUFFER_SIZE = int(5e5)  # replay buffer size
    BATCH_SIZE = 256        # minibatch size
    GAMMA = 0.99             # discount factor
    TAU = 1e-3               # for soft update of target parameters
    LR_ACTOR = 1e-3          # learning rate of the actor
    LR_CRITIC = 1e-3         # learning rate of the critic
    WEIGHT_DECAY = 0         # L2 weight decay

    device = "cpu"

```

4.4 Getting data about an environment

```

[4]: # get the default brain
    brain_name = env.brain_names[0]
    brain = env.brains[brain_name]

```

4.5 Creating class Actor that will act on behalf of the agent

```

[5]: def hidden_init(layer):
        fan_in = layer.weight.data.size()[0]
        lim = 1. / np.sqrt(fan_in)
        return (-lim, lim)

    class Actor(nn.Module):
        """Actor (Policy) Model."""

        def __init__(self, state_size, action_size, seed, fc1_units=128,
            →fc2_units=128):
            """Initialize parameters and build model.
            Params
            =====
                state_size (int): Dimension of each state

```

```

        action_size (int): Dimension of each action
        seed (int): Random seed
        fc1_units (int): Number of nodes in first hidden layer
        fc2_units (int): Number of nodes in second hidden layer
        """
        super(Actor, self).__init__()
        self.seed = torch.manual_seed(seed)
        self.fc1 = nn.Linear(state_size, fc1_units)
        self.fc2 = nn.Linear(fc1_units, fc2_units)
        self.fc3 = nn.Linear(fc2_units, action_size)
        self.reset_parameters()

    def reset_parameters(self):
        self.fc1.weight.data.uniform_(*hidden_init(self.fc1))
        self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
        self.fc3.weight.data.uniform_(-3e-3, 3e-3)

    def forward(self, state):
        """Build an actor (policy) network that maps states -> actions."""
        x = F.relu(self.fc1(state))
        x = F.relu(self.fc2(x))
        return F.tanh(self.fc3(x))

```

4.5 Creating class Critic which “criticize” our class Actor

```

[6]: class Critic(nn.Module):
        """Critic (Value) Model."""

        def __init__(self, state_size, action_size, seed, fcs1_units=128,
        ↪fc2_units=128):
            """Initialize parameters and build model.
            Params
            =====
            state_size (int): Dimension of each state
            action_size (int): Dimension of each action
            seed (int): Random seed
            fcs1_units (int): Number of nodes in the first hidden layer
            fc2_units (int): Number of nodes in the second hidden layer
            """
            super(Critic, self).__init__()
            self.seed = torch.manual_seed(seed)
            self.fcs1 = nn.Linear(state_size, fcs1_units)
            self.fc2 = nn.Linear(fcs1_units+action_size, fc2_units)
            self.fc3 = nn.Linear(fc2_units, 1)
            self.reset_parameters()

        def reset_parameters(self):

```

```

self.fc1.weight.data.uniform_(*hidden_init(self.fc1))
self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
self.fc3.weight.data.uniform_(-3e-3, 3e-3)

def forward(self, state, action):
    """Build a critic (value) network that maps (state, action) pairs -> Q-values."""
    xs = F.relu(self.fc1(state))
    x = torch.cat((xs, action), dim=1)
    x = F.relu(self.fc2(x))
    return self.fc3(x)

```

4.6 Creating class Agent that will represent our agent The most important component is a method learn:

Instantiating actors.

```

self.actor_local = Actor(state_size, action_size, random_seed).to(device)
self.actor_target = Actor(state_size, action_size, random_seed).to(device)
self.actor_optimizer = optim.Adam(self.actor_local.parameters(), lr=LR_ACTOR)

```

Instantiating critics.

```

self.critic_local = Critic(state_size, action_size, random_seed).to(device)
self.critic_target = Critic(state_size, action_size, random_seed).to(device)
self.critic_optimizer = optim.Adam(self.critic_local.parameters(), lr=LR_CRITIC, weight_decay=1e-4)

```

actions_next = self.actor_target(next_states) - getting target actions.

Q_targets_next = self.critic_target(next_states, actions_next) - getting target Q value.

Q_targets = rewards + (gamma * Q_targets_next * (1 - dones)) - getting updated target Q value.

Q_expected = self.critic_local(states, actions) - getting current Q value.

critic_loss = F.mse_loss(Q_expected, Q_targets) - decreasing difference between current Q value and updated target Q value.

actions_pred = self.actor_local(states) - getting current actions.

actor_loss = -self.critic_local(states, actions_pred).mean() - updating critic_local. Increasing Q values using gradient ascent.

Making target closer to local with coefficient TAU.

```
self.soft_update(self.critic_local, self.critic_target, TAU)
```

```
self.soft_update(self.actor_local, self.actor_target, TAU)
```

self.epsilon -= self.epsilon_decay - decreasing exploration exploitation rate.

```

[7]: class Agent():
    """Interacts with and learns from the environment."""

    def __init__(self, state_size, action_size, random_seed, epsilon=1.0,
        epsilon_decay=1e-6):

```

```

"""Initialize an Agent object.

Params
=====
    state_size (int): dimension of each state
    action_size (int): dimension of each action
    random_seed (int): random seed
"""

self.epsilon = epsilon
self.epsilon_decay = epsilon_decay

self.state_size = state_size
self.action_size = action_size
self.seed = random.seed(random_seed)

# Actor Network (w/ Target Network)
self.actor_local = Actor(state_size, action_size, random_seed).
→to(device)
self.actor_target = Actor(state_size, action_size, random_seed).
→to(device)
self.actor_optimizer = optim.Adam(self.actor_local.parameters(),
→lr=LR_ACTOR)

# Critic Network (w/ Target Network)
self.critic_local = Critic(state_size, action_size, random_seed).
→to(device)
self.critic_target = Critic(state_size, action_size, random_seed).
→to(device)
self.critic_optimizer = optim.Adam(self.critic_local.parameters(),
→lr=LR_CRITIC, weight_decay=WEIGHT_DECAY)

# Noise process
self.noise = OUNoise(action_size, random_seed)

# Replay memory
self.memory = ReplayBuffer(action_size, BUFFER_SIZE, BATCH_SIZE,
→random_seed)
# self.cache = ReplayBuffer(action_size, int(6e4), BATCH_SIZE,
→random_seed)

for target_param, param in zip(self.actor_target.parameters(), self.
→actor_local.parameters()):
    target_param.data.copy_(param.data)
for target_param, param in zip(self.critic_target.parameters(), self.
→critic_local.parameters()):
    target_param.data.copy_(param.data)

```



```

def step(self, state, action, reward, next_state, done, epoch):
    """Save experience in replay memory, and use random sample from buffer
    to learn."""
    # Save experience / reward
    self.memory.add(state, action, reward, next_state, done)
    # for states, actions, rewards, next_states, dones in zip(state,
    action, reward, next_state, done):
    self.memory.add(states, actions, rewards, next_states, dones)
    self.memory.add(state, action, reward, next_state, done)

    self.cache.add(states, actions, rewards, next_states, dones)

    # Learn, if enough samples are available in memory
    if len(self.memory) > BATCH_SIZE and epoch % 20 == 0:
        for num in range(18):
            experiences = self.memory.sample()
            self.learn(experiences, GAMMA)
        # for num in range(3):
        #     experiences = self.cache.sample()
        #     self.learn(experiences, GAMMA)

def act(self, state, add_noise=True):
    """Returns actions for given state as per current policy."""
    state = torch.from_numpy(state).float().to(device)
    self.actor_local.eval()
    with torch.no_grad():
        action = self.actor_local(state).cpu().data.numpy()
    self.actor_local.train()
    if add_noise:
        action += self.epsilon * self.noise.sample()
    return action

def reset(self):
    self.noise.reset()

def learn(self, experiences, gamma):
    """Update policy and value parameters using given batch of experience
    tuples.
    Q_targets = r + * critic_target(next_state, actor_target(next_state))
    where:
        actor_target(state) -> action
        critic_target(state, action) -> Q-value

```

Params

```

=====
    experiences (Tuple[torch.Tensor]): tuple of (s, a, r, s', done)
→ tuples
    gamma (float): discount factor
    """
    states, actions, rewards, next_states, dones = experiences

    # ----- update critic
→ ----- #
    # Get predicted next-state actions and Q values from target models
    actions_next = self.actor_target(next_states)
    Q_targets_next = self.critic_target(next_states, actions_next)
    # Compute Q targets for current states (y_i)
    Q_targets = rewards + (gamma * Q_targets_next * (1 - dones))
    # Compute critic loss
    Q_expected = self.critic_local(states, actions)
    critic_loss = F.mse_loss(Q_expected, Q_targets)
    # Minimize the loss

    self.critic_optimizer.zero_grad()
    critic_loss.backward()
    torch.nn.utils.clip_grad_norm_(self.critic_local.parameters(), 1)
    self.critic_optimizer.step()

    # ----- update actor
→ ----- #
    # Compute actor loss
    actions_pred = self.actor_local(states)
    actor_loss = -self.critic_local(states, actions_pred).mean()
    # Minimize the loss
    self.actor_optimizer.zero_grad()
    actor_loss.backward()
    self.actor_optimizer.step()

    # ----- update target networks
→ ----- #
    self.soft_update(self.critic_local, self.critic_target, TAU)
    self.soft_update(self.actor_local, self.actor_target, TAU)
    self.epsilon -= self.epsilon_decay
    self.noise.reset()

def soft_update(self, local_model, target_model, tau):
    """Soft update model parameters.
    _target = *_local + (1 - )*_target

```

Params

```

        =====
        local_model: PyTorch model (weights will be copied from)
        target_model: PyTorch model (weights will be copied to)
        tau (float): interpolation parameter
    """
    for target_param, local_param in zip(target_model.parameters(),
    ↪ local_model.parameters()):
        target_param.data.copy_(tau*local_param.data + (1.
    ↪ 0-tau)*target_param.data)

```

4.7 Noise One of the key components of DDPG. Ornstein-Uhlenbeck makes exploration more effective.

```

[8]: class OUNoise:
    """Ornstein-Uhlenbeck process."""

    def __init__(self, size, seed, mu=0., theta=0.15, sigma=0.05):
        """Initialize parameters and noise process."""
        self.mu = mu * np.ones(size)
        self.theta = theta
        self.sigma = sigma
        self.seed = random.seed(seed)
        self.reset()

    def reset(self):
        """Reset the internal state (= noise) to mean (mu)."""
        self.state = copy.copy(self.mu)

    def sample(self):
        """Update internal state and return it as a noise sample."""
        x = self.state
        dx = self.theta * (self.mu - x) + self.sigma * np.array([random.
    ↪ random() for i in range(len(x))])
        self.state = x + dx
        return self.state

```

4.8 Replay buffer

```

[9]: class ReplayBuffer:
    """Fixed-size buffer to store experience tuples."""

    def __init__(self, action_size, buffer_size, batch_size, seed):
        """Initialize a ReplayBuffer object.
        Params
        =====
        buffer_size (int): maximum size of buffer
        batch_size (int): size of each training batch

```

```

        """
        self.action_size = action_size
        self.memory = deque(maxlen=buffer_size)  # internal memory (deque)
        self.batch_size = batch_size
        self.experience = namedtuple("Experience", field_names=["state",
→"action", "reward", "next_state", "done"])
        self.seed = random.seed(seed)

    def add(self, state, action, reward, next_state, done):
        """Add a new experience to memory."""
        for num in range(20):
            e = self.experience(state[num], action[num], reward[num],
→next_state[num], done[num])
            self.memory.append(e)

    def sample(self):
        """Randomly sample a batch of experiences from memory."""

        experiences = random.choices(self.memory, k=self.batch_size)
        states = torch.from_numpy(np.vstack([e.state for e in experiences if e
→is not None])).float().to(device)
        actions = torch.from_numpy(np.vstack([e.action for e in experiences if
→e is not None])).float().to(device)
        rewards = torch.from_numpy(np.vstack([e.reward for e in experiences if
→e is not None])).float().to(device)
        next_states = torch.from_numpy(np.vstack([e.next_state for e in
→experiences if e is not None])).float().to(device)
        dones = torch.from_numpy(np.vstack([e.done for e in experiences if e is
→not None])).astype(np.uint8).float().to(device)

        return (states, actions, rewards, next_states, dones)

    def __len__(self):
        """Return the current size of internal memory."""
        return len(self.memory)

```

4.9 Instantiating class Agent

```
[10]: agent = Agent(state_size=33, action_size=4, random_seed=2)
```

5.0 Training loop

```
[11]: def ddpg(n_episodes=1000):
    print('Start time: ',datetime.datetime.now())
    all_agents = np.array([]).reshape(0,20)
    start_time = time.time()
    print('\rEP, Min, Max, Average, AV100, Time')
```

```

    for i_episode in range(1, n_episodes+1):
        env_info = env.reset(train_mode=True)[brain_name]           # reset the
→environment
        states = env_info.vector_observations                       # get the
→current state (for each agent)
        scores = np.zeros(20)                                     # initialize the score
→(for each agent)
        agent.reset()

        score_average = 0
        timestep = time.time()
        for t in range(100000000):
            actions = agent.act(states, add_noise=True)
            env_info = env.step(actions)[brain_name]               # send all
→actions to the environment
            next_states = env_info.vector_observations             # get next state
→(for each agent)
            rewards = env_info.rewards                             # get reward
→(for each agent)
            dones = env_info.local_done                            # see if episode
→finished
            agent.step(states, actions, rewards, next_states, dones,t)
            states = next_states                                   # roll over
→states to next time step
            scores += rewards                                     # update the
→score (for each agent)
            if np.any(dones):                                     # exit loop if
→episode finished
                break

        all_agents = np.concatenate((all_agents, scores.reshape(1,20)), axis=0)
        mean_100 = np.mean(all_agents[-101:-1], axis=0)
        print('{} , {:.2f}, {:.2f}, {:.2f}, {:.2f}, {:.2f}'\
            .format(str(i_episode).zfill(3), np.min(scores), np.max(scores),
→np.mean(scores), float(np.mean(mean_100)),
            time.time() - timestep), end="\n")
        if np.all(mean_100 > 30.0):
            end_time = time.time()
            print('\nSolved in {:d} episodes!\tAvg Score: {:.2f}, time: {}'.
→format(i_episode, float(np.mean(mean_100)),
            end_time-start_time))
            for num, scr in enumerate(mean_100):
                print('Agent {} average: {:.2f} '.format(num+1, scr))
            torch.save(agent.actor_local.state_dict(), 'checkpoint_actor.pth')
            torch.save(agent.critic_local.state_dict(), 'checkpoint_critic.pth')

```

```
print('End time: ',datetime.datetime.now())
return all_agents
```

5.1 Training

```
[12]: scores = ddpg()
```

Start time: 2019-08-12 21:10:10.629150

EP, Min, Max, Average, AV100, Time

```
C:\Users\andreiliphd\Anaconda3\envs\drlnd\lib\site-
packages\numpy\core\fromnumeric.py:3118: RuntimeWarning: Mean of empty slice.
    out=out, **kwargs)
C:\Users\andreiliphd\Anaconda3\envs\drlnd\lib\site-
packages\numpy\core\_methods.py:78: RuntimeWarning: invalid value encountered in
true_divide
    ret, rcount, out=ret, casting='unsafe', subok=False)
C:\Users\andreiliphd\Anaconda3\envs\drlnd\lib\site-
packages\ipykernel_launcher.py:31: RuntimeWarning: invalid value encountered in
greater
```

```
001, 0.00, 1.69, 0.47, nan, 12.37
002, 0.00, 0.21, 0.06, 0.47, 12.07
003, 0.00, 1.16, 0.17, 0.26, 12.21
004, 0.00, 0.65, 0.10, 0.23, 12.31
005, 0.00, 1.96, 0.28, 0.20, 12.38
006, 0.00, 0.90, 0.27, 0.22, 13.00
007, 0.00, 0.90, 0.38, 0.22, 13.01
008, 0.43, 2.27, 1.20, 0.25, 13.18
009, 0.72, 2.30, 1.24, 0.37, 13.35
010, 0.60, 2.59, 1.49, 0.46, 13.46
011, 0.61, 2.57, 1.42, 0.57, 13.47
012, 0.45, 4.85, 1.54, 0.64, 13.87
013, 0.00, 2.52, 1.28, 0.72, 14.03
014, 0.00, 3.81, 1.62, 0.76, 14.26
015, 0.16, 3.45, 1.69, 0.82, 14.52
016, 0.66, 4.14, 2.00, 0.88, 15.21
017, 0.47, 3.93, 2.18, 0.95, 15.48
018, 0.54, 3.93, 1.79, 1.02, 15.56
019, 1.11, 3.71, 2.51, 1.07, 15.83
020, 1.21, 4.15, 2.29, 1.14, 16.04
021, 0.95, 5.42, 3.38, 1.20, 16.13
022, 1.49, 4.41, 2.90, 1.30, 16.40
023, 0.81, 4.72, 2.73, 1.38, 16.37
024, 0.53, 6.13, 3.31, 1.43, 16.72
025, 0.35, 4.87, 2.52, 1.51, 17.22
026, 0.64, 6.78, 2.74, 1.55, 17.32
027, 1.14, 5.31, 3.21, 1.60, 17.36
```

028, 1.50, 7.55, 3.99, 1.66, 17.28
029, 0.65, 6.42, 3.20, 1.74, 17.57
030, 1.09, 8.87, 4.69, 1.79, 17.46
031, 2.03, 6.89, 4.16, 1.89, 17.56
032, 1.74, 6.62, 3.86, 1.96, 17.87
033, 2.63, 7.86, 5.08, 2.02, 17.83
034, 2.01, 8.74, 4.97, 2.11, 17.82
035, 2.02, 9.46, 6.21, 2.20, 17.70
036, 3.59, 10.80, 7.10, 2.31, 17.68
037, 1.44, 18.65, 7.28, 2.45, 17.81
038, 3.05, 11.94, 8.19, 2.58, 17.65
039, 2.60, 10.95, 6.57, 2.72, 17.85
040, 3.78, 15.61, 9.04, 2.82, 17.71
041, 3.50, 15.97, 8.43, 2.98, 17.83
042, 4.01, 16.52, 9.49, 3.11, 17.63
043, 2.19, 16.30, 10.47, 3.26, 17.58
044, 4.18, 15.49, 9.50, 3.43, 17.72
045, 3.91, 15.52, 10.84, 3.57, 17.82
046, 3.11, 20.63, 10.36, 3.73, 18.45
047, 7.19, 21.93, 12.95, 3.87, 17.63
048, 5.61, 18.83, 12.83, 4.07, 17.43
049, 6.00, 20.54, 12.99, 4.25, 17.75
050, 5.83, 17.45, 12.69, 4.43, 17.72
051, 9.59, 22.73, 15.28, 4.59, 17.43
052, 8.55, 22.84, 15.20, 4.80, 17.35
053, 8.35, 24.88, 16.45, 5.00, 17.48
054, 10.46, 22.23, 16.54, 5.22, 17.40
055, 10.68, 23.67, 18.08, 5.43, 17.61
056, 10.74, 23.39, 17.61, 5.66, 17.58
057, 12.37, 23.63, 17.85, 5.87, 17.76
058, 14.05, 23.23, 18.30, 6.08, 17.56
059, 14.89, 27.22, 18.22, 6.29, 17.55
060, 9.20, 23.53, 16.57, 6.50, 17.64
061, 12.12, 26.17, 20.31, 6.66, 17.76
062, 9.82, 27.78, 19.36, 6.89, 17.80
063, 11.09, 26.26, 19.62, 7.09, 17.73
064, 13.96, 25.83, 19.56, 7.29, 17.71
065, 11.82, 25.19, 18.82, 7.48, 17.74
066, 14.86, 30.43, 20.97, 7.65, 17.71
067, 11.29, 27.56, 18.65, 7.86, 17.78
068, 10.24, 25.60, 19.26, 8.02, 17.83
069, 11.39, 28.27, 20.14, 8.18, 17.63
070, 10.81, 28.18, 20.71, 8.36, 17.74
071, 13.92, 35.56, 22.13, 8.53, 17.97
072, 14.99, 28.90, 22.30, 8.72, 18.20
073, 9.39, 28.36, 21.65, 8.91, 17.88
074, 15.89, 27.23, 24.06, 9.09, 17.85
075, 9.31, 27.16, 21.96, 9.29, 17.82

076, 12.31, 33.18, 22.40, 9.46, 17.62
077, 7.84, 29.90, 19.48, 9.63, 17.65
078, 6.76, 31.16, 23.96, 9.76, 17.54
079, 8.41, 36.72, 23.50, 9.94, 17.52
080, 16.67, 30.35, 24.22, 10.11, 17.42
081, 13.67, 31.56, 24.84, 10.29, 17.48
082, 10.46, 32.97, 22.46, 10.47, 17.59
083, 12.17, 30.28, 22.04, 10.61, 17.69
084, 15.68, 28.89, 23.73, 10.75, 17.70
085, 15.41, 37.72, 24.55, 10.90, 17.49
086, 14.81, 29.72, 23.56, 11.06, 17.46
087, 18.26, 30.81, 25.26, 11.21, 17.49
088, 18.35, 32.07, 25.39, 11.37, 17.27
089, 20.66, 32.81, 26.34, 11.53, 17.58
090, 23.38, 36.71, 30.11, 11.70, 17.69
091, 17.01, 34.01, 29.07, 11.90, 17.47
092, 14.44, 35.39, 27.31, 12.09, 17.71
093, 16.32, 36.04, 29.12, 12.26, 17.48
094, 20.17, 35.51, 27.12, 12.44, 17.37
095, 20.47, 35.41, 28.76, 12.59, 17.32
096, 18.79, 35.24, 28.57, 12.76, 17.12
097, 20.61, 35.82, 27.94, 12.93, 17.33
098, 20.45, 35.18, 29.86, 13.08, 17.10
099, 16.72, 35.83, 28.90, 13.25, 17.15
100, 18.94, 36.16, 29.58, 13.41, 17.09
101, 17.49, 34.28, 27.47, 13.57, 17.13
102, 20.54, 39.04, 28.49, 13.84, 17.25
103, 26.74, 39.62, 31.55, 14.13, 16.95
104, 25.51, 38.98, 30.76, 14.44, 17.45
105, 25.22, 37.19, 30.68, 14.75, 17.36
106, 16.84, 36.14, 30.74, 15.05, 17.57
107, 18.66, 37.82, 27.59, 15.36, 17.35
108, 15.80, 37.80, 30.88, 15.63, 17.43
109, 15.92, 36.85, 28.09, 15.93, 17.59
110, 18.08, 37.15, 28.29, 16.19, 17.58
111, 18.66, 37.15, 28.67, 16.46, 17.75
112, 23.25, 37.02, 31.62, 16.74, 17.46
113, 21.29, 39.07, 34.32, 17.04, 17.53
114, 21.06, 39.54, 30.60, 17.37, 17.66
115, 25.24, 38.68, 32.54, 17.66, 17.44
116, 21.18, 37.86, 31.32, 17.96, 17.40
117, 20.61, 36.48, 30.04, 18.26, 17.77
118, 17.24, 36.48, 31.48, 18.54, 17.64
119, 25.56, 38.76, 31.85, 18.83, 17.29
120, 23.32, 39.53, 33.89, 19.13, 17.54
121, 21.67, 39.59, 33.53, 19.44, 17.40
122, 25.76, 39.05, 34.53, 19.74, 17.38
123, 15.50, 39.23, 33.93, 20.06, 17.55

124, 29.37, 39.39, 35.19, 20.37, 17.25
125, 21.09, 39.60, 34.61, 20.69, 17.54
126, 29.32, 39.49, 36.37, 21.01, 17.23
127, 35.35, 39.31, 37.83, 21.35, 17.15
128, 20.84, 39.56, 35.36, 21.69, 17.20
129, 26.26, 39.66, 36.93, 22.01, 17.07
130, 18.73, 39.63, 36.22, 22.35, 17.26
131, 28.73, 39.61, 37.78, 22.66, 17.24
132, 33.09, 39.67, 37.53, 23.00, 17.31
133, 25.13, 39.66, 36.66, 23.33, 17.31
134, 26.09, 39.38, 37.31, 23.65, 17.31
135, 28.21, 39.65, 36.85, 23.97, 17.41
136, 32.54, 39.62, 36.98, 24.28, 17.39
137, 30.32, 39.58, 35.67, 24.58, 17.56
138, 29.89, 39.66, 36.02, 24.86, 17.45
139, 29.52, 39.65, 35.36, 25.14, 17.51
140, 24.31, 39.57, 35.67, 25.43, 17.55
141, 31.41, 39.17, 36.89, 25.69, 17.51
142, 32.77, 39.54, 37.10, 25.98, 17.59
143, 32.01, 39.10, 36.81, 26.25, 17.36
144, 23.34, 39.30, 35.69, 26.52, 17.32
145, 31.75, 38.87, 36.44, 26.78, 17.26
146, 33.20, 39.02, 36.49, 27.04, 17.38
147, 33.21, 38.91, 36.36, 27.30, 17.48
148, 31.44, 39.14, 36.26, 27.53, 17.23
149, 28.54, 38.90, 35.77, 27.77, 17.38
150, 29.19, 38.85, 35.75, 27.99, 17.25
151, 33.54, 39.65, 37.39, 28.22, 17.37
152, 21.88, 39.51, 35.92, 28.45, 17.33
153, 29.00, 39.37, 36.58, 28.65, 17.07
154, 30.87, 39.61, 37.55, 28.85, 17.18
155, 29.71, 39.55, 37.33, 29.06, 17.17
156, 20.55, 39.38, 34.47, 29.26, 17.29
157, 21.84, 39.45, 35.95, 29.42, 17.24
158, 35.89, 39.53, 38.22, 29.61, 17.77
159, 33.55, 39.66, 37.60, 29.81, 17.78
160, 32.79, 39.13, 37.11, 30.00, 17.52
161, 32.65, 39.49, 37.12, 30.20, 17.71
162, 25.02, 39.63, 35.26, 30.37, 17.59
163, 25.97, 37.57, 35.04, 30.53, 17.48
164, 28.69, 39.43, 36.25, 30.69, 17.53
165, 30.66, 39.64, 36.21, 30.85, 17.39
166, 30.38, 39.47, 36.66, 31.03, 17.53

Solved in 166 episodes! Avg Score: 31.03, time: 2829.3051903247833

Agent 1 average: 30.92

Agent 2 average: 31.59

Agent 3 average: 31.73

```
Agent 4 average: 31.03
Agent 5 average: 30.52
Agent 6 average: 31.17
Agent 7 average: 31.70
Agent 8 average: 30.96
Agent 9 average: 30.93
Agent 10 average: 30.56
Agent 11 average: 31.01
Agent 12 average: 30.85
Agent 13 average: 31.57
Agent 14 average: 30.83
Agent 15 average: 30.92
Agent 16 average: 30.62
Agent 17 average: 30.15
Agent 18 average: 30.71
Agent 19 average: 31.66
Agent 20 average: 31.07
End time: 2019-08-12 21:57:19.940336
```


```
[13]: env.close()
```

6. Visualizing rewards accross 20 workers

```
[29]: plt.plot(np.arange(1, len(scores)+1), scores)
      plt.ylabel('Rewards')
      plt.xlabel('Episode #')
```

WARNING:matplotlib.legend:No handles with labels found to put in legend.

```
[29]: <matplotlib.legend.Legend at 0x267ac514c18>
```



1.0.6 Improvements

1. Batch normalization could be added to improve numerical stability.
2. Sigma coefficient might be tuned for a faster convergence.
3. Neural network might be tuned for a better performance. Too many layers and too many neurons lead to a slower convergence.

DDPG looks for me very fragile. One step right or left and an implementation should be retuned.

[]: